

January 2025 CSE 314: Operating System Sessional

Assignment 4: Advanced Memory Management for xv6

Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology

June 2025

I welcome you all to one of the most fascinating assignments of your undergraduate life: implementing advanced memory management in xv6. I hope you enjoy working on it and find it rewarding once you complete it.

1 Introduction

xv6 implements a very basic memory management system. In this assignment, you will enhance it to a great extent; namely, you will add COW and Paging.

2 Paging

As the memory is normally very small compared to the need for a process. Paging operation puts some pages to persistent storage (disk) when memory gets full. When a swapped-out page is needed again, it is brought back into memory. It makes available virtual memory much larger than physical memory. The decision on which page to swap out is determined by the page replacement algorithm in use. Follow these steps to implement paging mechanism in xv6 operating system.

1. Download and apply the following patch file.

[swap.patch](#)

This patch file will help you swap out and in pages to and from disk respectively. Understand how it does its job. Update any other code that this code depends on to work.

It contains a struct named `swap`. This structure contains the metadata to retrieve a page that has been swapped out to disk. When a page is *swapped out*, it saves the block no. to the blocks that store that page. A swapped out page can be *swapped in* using the swap struct that was created when that page was swapped out.

It implements the following functions:

- `void swapinit(void)`: Initializes necessary variables for allocating and freeing swap structs.

- `struct swap* swapalloc(void)`: Allocates a swap struct and returns a pointer to that struct.
- `void swapfree(struct swap*)`: Frees a given swap struct that can be reused during some future `swapalloc()`.
- `void swapout(struct swap *dst_sp, char *src_pa)`: Writes the page `src_pa` to disk and saves the block no.s to `dst_sp`.
- `void swapin(char *dst_pa, struct swap *src_sp)`: Reads a page into `dst_pa` that has been previously swapped out (by calling `swapout()`) with `src_sp` as argument.

Write these two functions to do the swap operations.

- **Swap out**: Allocate a swap struct using `swapalloc`. Then use `swapout` to write the page to disk. The swap struct should be saved somewhere so that you can swap in this page when needed.
- **Swap in**: First find the swap struct that was created when the page was swapped out. Then use `swapin` with this struct to copy the data from disk.

Note: `swapout` and `swapin` does neither allocate or deallocate anything.

2. You need to keep all the live pages in a data structure so that you can decide which page to swap out. Use a linked list for this task. (A simple array will also work but you have to use linked list to get full marks). Check if you can keep track of all pages being used across different processes.

Test code 1: This is a good place to test your implementation. Write a user code that uses some number of pages provided by the command line. Make sure the user code takes some time to execute. Write a system-call that prints the number of **live** pages being used by different processes. Now run multiple instances of the user code from shell and use the system-call to check if the counts match.

3. At any time, you should have no more than `MAXPHYPAGES` (=50) pages being used by all user processes in total. Implement a simple page replacement algorithm (I recommend FIFO). Make sure you correctly free the swap structs and physical pages.

Test code 2: This is an important place to test your implementation. Check the total number of live pages and how it changes after swapping in/out. You should test after each small update of the code and check if it works as expected (for example, first check if you can swap out pages and reduce the number of live pages; then check if you can swap in a page when required.)

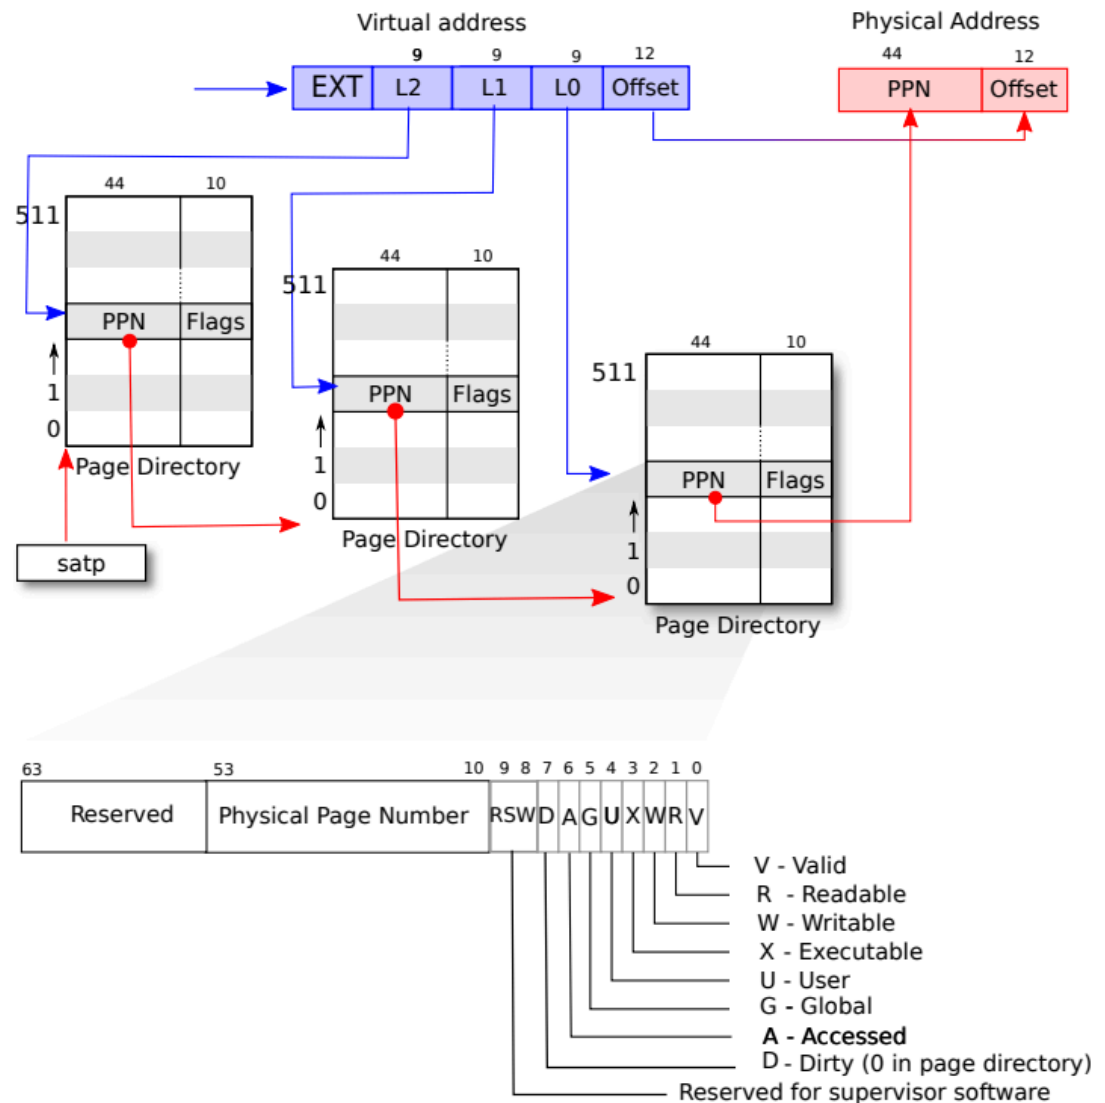


Figure 3.2: RISC-V page table hardware.

4. Your simple implementation is not likely to work in the presence of `fork()`. You need to keep the number of references to each swap struct so that you can garbage collect. Update the swap struct accordingly.

Test code 3: Write appropriate test codes here. The user code should include some number of forks. Make sure the parent process is big enough so that some page is already swapped out before `fork` is called, then the child process has to access a swapped out page from the parent process.

3 Bonus

Implement CoW (Copy on Write on top of paging mechanism). For more information, Check the next section.

CoW (Copy-on-Write)

CoW or Copy-On-Write is an improvement on regular memory allocation of `fork()` system-call. In xv6, this system-call directly copies all the *user-space memory* of the parent process to the child. Now, if a page is only ever used to read (or execute), then it creates unnecessary copy. COW tries to minimize this by deferring the copy until a write operation is done on a page. It sets up the pagetable in the child process such that all the PTEs for *user memory* point to the parent's physical pages and clears the *write flags* of all PTEs for *user memory* in both child and parent. When a write operation is done on any non-writable page, the CPU forces a page fault. Kernel keeps track of which pages are COW and when these faults occur for a COW page, it allocates a new page for the faulting process, copies the data from the COW page to it and updates the corresponding PTE such that it points to this new page with the *write flag* set. Now, the write operation can carry on. Follow these steps to implement CoW on top of paging mechanism on xv6.

1. Find out and understand the function that creates pagetable of the child process during `fork()` system-call. [Hint: it is `uvmcopy()` in `kernel/vm.c`.]
2. You need to handle the page fault that occurs while writing in a COW page. When any type of trap occurs from user code (page fault belongs to this type), it calls `usertrap()` in `kernel/trap.c`. Understand how the trap for system-call is handled in this code. Also, check what happens for any other type of trap. [Don't get bothered with `which_dev.devintr()` returns 0 when the CPU itself generates the trap (like system-call, page fault).]
3. Update the function you found in step 1 so that no new page is allocated, and the PTEs in the child process' pagetable points to the parent's physical pages. Also, make sure that the *write flag* is clear for both parent and child processes. [Many macros and definitions related to page tables can be found at the end of `kernel/riscv.h`]
4. Update the `usertrap()` function in `kernel/trap.c` so that write operation on a COW page creates a writable copied page for the faulting process. So, you need to keep track of if a page is COW. PTEs in riscv have 10 flags of which 8 are used by hardware as in Figure 3.2 and 2 other reserved for software (RSW) (although xv6 provides definition for first 4 bits as given in `kernel/riscv.h`). You can use any of the RSW bits. [Original `uvmcopy()` you modified already implements this copying process (albeit for all pages). Take inspiration from it.]

5. You may think it is done. Unfortunately, no! If you have carefully followed the process, you can see once a page is set non-writable, it is never set writable. Now, if both parent and child write on the same COW page, the non-writable COW page is lost (as both parent and child created its copy), i.e., memory leak by kernel! Also, a COW page can be referenced by many processes (depending on number of forks). The only way to solve it is to keep a “reference count” for each page so that a “kalloc”ed page can be garbage collected. You can choose various schemes to keep the reference counts for all pages. You may use a fixed size array where you can index by the page's physical address divided by page size.
6. Now, we should be done. Or are we? Unfortunately, there is still a problem. Up to this point, we assumed that the write operation to a COW page is only done in user mode. But it is possible that a write to a COW page is done in kernel mode (for example, `copyout()` in `kernel/vm.c`). A page fault in kernel mode generates a kernel trap which is handled by `kerneltrap()` in `kernel/trap.c` as opposed to `usertrap()`. You can solve this issue by changing `copyout()` and any other functions like it, or by changing `kerneltrap()`.
[Hint:
 1. *`copyout()` gets called for pipe and file I/O. So, your test code should do some file I/O or pipe operation before and after `fork()`.*
 2. *You can also write a new system-call that uses `copyout()`. Your test code will call this system-call before and after `fork()`.]*
7. Now, update your codes so that it supports COW on top of paging. You need a protocol for what to do with COW pages. Here you have two options: (1) not to swap a COW page, or (2) in the data structure in step 2, keep all (process, virtual address) pairs corresponding to each physical page *[Note: 1. This needs another linked list. 2. You need to update this list during `fork`, `sbrk` (with negative value) and `exit`.]* and when swapping a page, reset the valid bit for corresponding process' virtual addresses.
Option 1 is good enough for this offline.

4 Important instructions

- Don't start implementation right away. First understand what actually happens in xv6. During implementation, try to test after each small change and make sure everything runs as expected (maybe it is supposed to panic).
- If you cannot find where a kernel function is implemented, check out `kernel/defs.h`.
- **Keep the testing codes. The codes carry marks. All green colored texts ask for such code.**

- If you Implement the bonus task CoW, you can implement CoW and Paging independently first for your convenience if you want. Keep their patch files. Submit these two patch files if you cannot do a combined implementation (You will get partial bonus marks).
- If you have one patch file, submit that file renamed by your student ID. Otherwise, put the patch files in a folder named by your student ID (ex: 2105200), zip it and submit. Make sure all the necessary codes are in the patch files so that it does not have any external dependency.
- **Do not copy. Any proof of copy will result in -100%.**
- **Submission Deadline: 11.55 PM, 19th July 2025**
- For any queries, you can contact me at amsrumi@gmail.com

5 Marks distribution

Task	Sub-task	Marks
Implementing Paging mechanism	<i>Detailed distribution will be declared later</i>	100
Bonus	Implementing CoW	15
Total		100+15